

Exercise set 2

your candidate number

High-dimensional regression

Background reading. Delete this section before submitting Wikipedia: Singular value decomposition. Read the introduction, Example, and Pseudoinverse sections.

Wikipedia: Pseudoinverse. Read the introduction and sections on Projectors, Examples, Linearly independent rows, and Applications.

ISLR: See section 12.5.2 for some example use of the `svd` function.

Delete from here to the previous line about deletion.

```
set.seed(1) # change this to some other number
# Generate a matrix of predictors X
# with 3 rows, 10 columns, and i.i.d. normal entries
p <- 10
n <- 3
X <- matrix(rnorm(n * p), nrow = n)
# Create a sparse coefficient vector beta
# with only 1 or 2 nonzero entries
beta <- rep(0, p)
beta[1] <- 1
# Compute outcome y from the noise-free linear model
y <- X %*% beta
```

Generate data (3 points)

Compute pseudoinverse (3 points) Use output from the `svd` function to compute a right pseudoinverse of X . (Note: if you use a source aside from the Wikipedia articles above to figure out how to do this you should cite your source and include a link if it's a website)

```
S <- svd(X)
X_pseudoinv <- S$v %*% diag(1/S$d) %*% t(S$u)
```

```
X %*% X_pseudoinv
```

Verify right-inverse property (3 points)

```
##           [,1]           [,2]           [,3]
## [1,]  1.000000e+00 -1.662692e-16 -1.289275e-16
## [2,] -3.247361e-16  1.000000e+00  1.401532e-16
## [3,] -1.886456e-16  3.672968e-16  1.000000e+00
```

Explanation: (1 point) Expected to be the identity matrix, and the diagonal entries are numerically close to 1 and off diagonal close to zero, i.e. the matrix is numerically close to the identity matrix.

Note: delete this comment after reading. If you are unable to compute the pseudoinverse using the `svd` function, you can increase the sample size to $p+1$ and use OLS to estimate β instead to receive partial credit.

```
beta_hat <- X_pseudoinv %*% y
cbind(beta, beta_hat) |> kable()
```

Compare estimated β to true β (3 points)

beta	
1	0.2152325
0	-0.1195434
0	-0.1159438
0	0.0664787
0	-0.1588557
0	-0.1190181
0	-0.2192990
0	0.1476748
0	-0.0767600
0	0.1468402

Explanation: (1 point) Expected they are not equal, because we do not have enough data to exactly solve for or estimate β

```
y_hat <- X %*% beta_hat
mean((y - y_hat)^2)
```

Compute MSE for predicting y (3 points)

```
## [1] 1.848893e-31
```

Explanation: (1 point) Expected to be numerically close to zero, because perfect prediction is possible when $p > n$.

```
X_test <- matrix(rnorm(n * p), nrow = n)
y_test <- X_test %*% beta
y_test_hat <- X_test %*% X_pseudoinv %*% y
mean((y_test_hat - y_test)^2)
```

Generate a new sample of test data and compute the (in-distribution) test MSE (3 points)

```
## [1] 0.7142776
```

Explanation: (1 point) Expected that this would be larger than zero, because test error is generally higher than training error.

Use penalized regression to estimate β (4 points) Compute the ridge and lasso estimates using $\lambda = 0.1$, and compare these with the estimate from using `svd`. You may want to read the documentation for `?glmnet` and `?coef.glmnet`

```
beta_ridge <- coef(glmnet(X, y, intercept = FALSE, alpha = 0, lambda = .1))
beta_lasso <- coef(glmnet(X, y, intercept = FALSE, lambda = .1))
```

```
# Comparison
```

```
cbind(beta, beta_lasso[-1], beta_ridge[-1], beta_hat) |>  
  kable()
```

beta				
1	0.7271206	0.2075285	0.2152325	
0	0.0000000	-0.0000405	-0.1195434	
0	0.0000000	0.4483238	-0.1159438	
0	0.0000000	0.1004163	0.0664787	
0	0.0000000	-0.0545385	-0.1588557	
0	0.0000000	-0.1072776	-0.1190181	
0	-0.1492376	-0.8173406	-0.2192990	
0	0.0000000	0.0249320	0.1476748	
0	0.0000000	-0.0941434	-0.0767600	
0	0.0000000	0.0019428	0.1468402	

Explanation: (1 point) Expected lasso solution to be sparse and ridge solution to not be sparse.

```
y_test_hat <- X_test %*% beta_lasso[-1]  
mean((y_test_hat - y_test)^2)
```

Compute test MSE using penalized regression estimates (3 points)

```
## [1] 0.05443351
```

```
y_test_hat <- X_test %*% beta_ridge[-1]  
mean((y_test_hat - y_test)^2)
```

```
## [1] 0.7803853
```

Explanation: (1 point) Expected or not, the lasso estimate is closest to the true beta and the test error is lower for the lasso compared to the other methods.