

ST455: Reinforcement Learning

Lecture 8: Policy Gradient Methods

Chengchun Shi

Recap

- We have covered **dynamic programming**, **Monte Carlo** methods (Lecture 3) and **temporal-difference learning** (Lectures 4 – 7). All these methods are **value-based** methods that derive the optimal policy by computing a value function.
- Today's lecture will adopt a difference approach that directly **searches** the optimal policy within a restricted policy class.

Lecture Outline

1. Policy Gradient Method Introduction
2. Policy Gradient Theorem
3. REINFORCE and Actor Critic Algorithms
4. Advantage Actor-Critic (A2C)

Lecture Outline

Policy Function Approximation

		No	Yes
Value Function Approximation	No	Value-based (tabular)	REINFORCE
	Yes	Value-based	Actor-Critic

- **Value-based**
 - Tabular (Lectures 3 & 4)
 - Function approx (Lectures 5 & 7)
- **REINFORCE**
 - No value function
 - Learn policy
- **Actor-critic**
 - Learn value
 - Learn policy
- **Advantage actor-critic**
 - Variance reduction

Lecture Outline

- 1. Policy Gradient Method Introduction**
2. Policy Gradient Theorem
3. REINFORCE and Actor Critic Algorithms
4. Advantage Actor-Critic (A2C)

Policy We Studied So Far

- Greedy policy:

$$\pi^{\text{opt}}(\textcolor{blue}{s}) = \arg \max_{\textcolor{red}{a}} Q^{\pi^{\text{opt}}}(\textcolor{blue}{s}, \textcolor{red}{a})$$

- ϵ -Greedy policy:


$$\begin{cases} \pi^{\text{opt}}(\textcolor{blue}{s}), & \text{with probability } \mathbf{1} - \epsilon \\ \text{random action,} & \text{with probability } \epsilon. \end{cases}$$

- **Value-based methods:** Policy Iteration, Value Iteration, SARSA, Q-Learning, etc.



Value-based v.s. Policy Gradient Methods

- **Value-based methods:** derive π^{opt} by learning an optimal Q-function (with or without function approximation)
- **Policy gradient methods:** search π^{opt} within a restricted function class (e.g., linear, neural networks) that maximizes the value

Value-based v.s. Policy Gradient Methods (Cont'd)



Value-based Methods



Policy Gradient Methods

Example: Linear Function Approximation

- Linear approximation of features $\phi(\mathbf{s}, \mathbf{a})$
- State-action value function approximation

$$Q(\mathbf{s}, \mathbf{a}; \theta) = \phi^\top(\mathbf{s}, \mathbf{a})\theta$$

- Policy function approximation


$$\pi(\mathbf{s}, \mathbf{a}; \theta) = \frac{\exp(\phi^\top(\mathbf{s}, \mathbf{a})\theta)}{\sum_{\mathbf{a}'} \exp(\phi^\top(\mathbf{s}, \mathbf{a}')\theta)}$$

$\phi^\top(\mathbf{s}, \mathbf{a})\theta$ similar to the preference score in the gradient based methods in HW1

Value-based v.s. Policy Gradient Methods (Cont'd)

- **Pros** of policy gradient methods:
 1. Suitable for learning general **stochastic** policies (value-based methods mainly designed for deterministic policies)
 2. More **robust** to model misspecification
 3. Scalable for **high-dimensional** or **continuous** action spaces (SARSA, Q-learning mainly designed for discrete action space)
- **Cons** of policy gradient methods:
 1. Convergence to local minima
 2. May have large variance


Example I: Advantage of Stochastic Policy



- **Reward** is -1 per step
- **Action:** left or right
- **G:** **terminal** state
- Second state: actions are **reversed**
- Linear combination of features
- $\phi(\mathbf{s}, \text{left}) = [1, 0]$ for any \mathbf{s}
- $\phi(\mathbf{s}, \text{right}) = [0, 1]$ for any \mathbf{s}

- π^{opt} : right \rightarrow left \rightarrow right
- Under the function approximation, implements the same policy at each state

Example I: Advantage of Stochastic Policy (Cont'd)



- **Reward** is -1 per step
- **Action**: left or right
- **G**: **terminal** state
- Second state: actions are **reversed**
- Linear combination of features
- $\phi(\mathbf{s}, \text{left}) = [1, 0]$ for any \mathbf{s}
- $\phi(\mathbf{s}, \text{right}) = [0, 1]$ for any \mathbf{s}

- Value-based method: $\arg \max_{\mathbf{a}} \phi^{\top}(\mathbf{s}, \mathbf{a})\boldsymbol{\theta} = \text{left}\mathbb{I}(\theta_1 > \theta_2) + \text{right}\mathbb{I}(\theta_1 \leq \theta_2)$
 - Deterministic: either always move to the left, or always move to the right
- Policy gradient method: $\pi(\text{right}|\mathbf{s}) = \exp(\theta_2)/[\exp(\theta_1) + \exp(\theta_2)]$
 - Stochastic: move to the right at each step with certain probability

Example II: Robustness of Policy Gradient Method

- Q-function is more **difficult** to model compared to the optimal policy
- Example: optimal Q-function: $Q^{\pi^{\text{opt}}}(\mathbf{s}, \mathbf{a}) = g(\phi^\top(\mathbf{s}, \mathbf{a})\theta^*)$ for some monotonically increasing function $g : \mathbb{R} \rightarrow \mathbb{R}$
- When g is not **identity** function, value-based method misspecifies Q-function model

$$g(\phi^\top(\mathbf{s}, \mathbf{a})\theta^*) \neq \phi^\top(\mathbf{s}, \mathbf{a})\theta$$

- However, since g is a monotonically increasing function

$$\pi^{\text{opt}}(\mathbf{s}) = \arg \max_{\mathbf{a}} g(\phi^\top(\mathbf{s}, \mathbf{a})\theta^*) = \arg \max_{\mathbf{a}} \phi^\top(\mathbf{s}, \mathbf{a})\theta^*$$

- Policy gradient methods correctly identifies the optimal policy

$$\frac{\exp(\phi^\top(\mathbf{s}, \mathbf{a})\theta)}{\sum_{\mathbf{a}'} \exp(\phi^\top(\mathbf{s}, \mathbf{a}')\theta)} \rightarrow \mathbb{I}(\mathbf{a} = \pi^{\text{opt}}(\mathbf{s}))$$

when $\theta = k\theta^*$ and $k \rightarrow \infty$

Policy Objective Functions

- Average rewards:

$$J(\theta) = \lim_{T \rightarrow \infty} \frac{1}{T} \mathbb{E}^{\pi(\bullet; \theta)} \left[\sum_{t=0}^{T-1} R_t \right] = \sum_{s, a} \nu^{\pi(\bullet; \theta)}(s) \pi(s, a; \theta) \mathcal{R}_s^a$$

where $\mathcal{R}_s^a = \mathbb{E}(R_t | A_t = a, S_t = s)$

- For each π , the states $\{S_t\}_t$ forms a time-homogeneous Markov chain
- $\nu^{\pi(\bullet; \theta)}$ the stationary distribution of $\{S_t\}_t$ under $\pi(\bullet; \theta)$

Policy Objective Functions (Cont'd)

- Discounted rewards: given a discounted factor $\gamma \in [0, 1]$ and initial state distribution ν , maximize the expected discounted rewards:

$$J(\theta) = \mathbb{E}^{\pi(\bullet; \theta)} \left[\sum_{t=0}^{\infty} \gamma^t R_t \right],$$

or equivalently,

$$J(\theta) = \sum_{\mathbf{s}} \nu(\mathbf{s}) V^{\pi(\bullet; \theta)}(\mathbf{s})$$

- If $\gamma = 1$, the task is assumed to be episodic

Lecture Outline

1. Policy Gradient Method Introduction
- 2. Policy Gradient Theorem**
3. REINFORCE and Actor Critic Algorithms
4. Advantage Actor-Critic (A2C)

Policy Gradient

- **Objective:** identify the maximizer of $J(\theta)$
- **Method:** apply (stochastic) gradient ascent algorithm to update θ (gradient descent to minimize $-J(\theta)$)

$$\theta_{t+1} = \theta_t + \alpha_t \nabla_{\theta} J(\theta_t)$$

Need to calculate the gradient $\nabla_{\theta} J(\theta)$!

Policy Gradient Theorem

Theorem

For any differentiable policy $\pi(\mathbf{s}, \mathbf{a}; \theta)$ with respect to parameter θ , the policy gradient for average reward and discounted expected rewards objective is

$$\nabla_{\theta} J(\theta) = \sum_{\mathbf{s}, \mathbf{a}} \mu^{\pi(\bullet; \theta)}(\mathbf{s}, \mathbf{a}) \nabla_{\theta} \log(\pi(\mathbf{s}, \mathbf{a}; \theta)) Q^{\pi(\bullet; \theta)}(\mathbf{s}, \mathbf{a})$$

- For average reward objective:
 $\mu^{\pi(\bullet; \theta)}$ is the stationary distribution of $\{(\mathbf{S}_t, \mathbf{A}_t)\}_t$ under $\pi(\bullet; \theta)$
- For discounted expected rewards objective:

$$\mu^{\pi(\bullet; \theta)}(\mathbf{s}, \mathbf{a}) = \sum_{t \geq 0} \gamma^t \Pr^{\pi(\bullet; \theta)}(\mathbf{S}_t = \mathbf{s}, \mathbf{A}_t = \mathbf{a})$$

Discounted state-action visitation probability

Policy Gradient Theorem (Cont'd)

Theorem

For any differentiable policy $\pi(\mathbf{s}, \mathbf{a}; \theta)$ with respect to parameter θ , the policy gradient for average reward and discounted expected rewards objective is

$$\nabla_{\theta} J(\theta) = \sum_{\mathbf{s}, \mathbf{a}} \mu^{\pi(\cdot; \theta)}(\mathbf{s}, \mathbf{a}) \nabla_{\theta} \log(\pi(\mathbf{s}, \mathbf{a}; \theta)) Q^{\pi(\cdot; \theta)}(\mathbf{s}, \mathbf{a})$$

- For average reward objective:

$$Q^{\pi}(\mathbf{s}, \mathbf{a}) = \mathbb{E}^{\pi} \left[\sum_{t \geq 0} (\mathbf{R}_t - J(\theta)) \mid \mathbf{S}_0 = \mathbf{s}, \mathbf{A}_0 = \mathbf{a} \right]$$

- For discounted expected rewards objective: Q-function defined as usual.
- Proof given in the appendix

Policy Score

- For any state-action pair (\mathbf{s}, \mathbf{a}), the term

$$\nabla_{\theta} \log(\pi(\mathbf{s}, \mathbf{a}; \theta))$$

is referred as the **policy score**

Example 1: Softmax Policy Gradient

- State-action pairs weighted by linear combination of features

$$\pi(\mathbf{s}, \mathbf{a}; \theta) = \frac{\exp(\phi^\top(\mathbf{s}, \mathbf{a})\theta)}{\sum_{\mathbf{a}'} \exp(\phi^\top(\mathbf{s}, \mathbf{a}')\theta)}$$

- The score function

$$\nabla_{\theta} \log \pi(\mathbf{s}, \mathbf{a}; \theta) = \phi(\mathbf{s}, \mathbf{a}) - \frac{\sum_{\mathbf{a}'} \phi(\mathbf{s}, \mathbf{a}') \exp(\phi^\top(\mathbf{s}, \mathbf{a}')\theta)}{\sum_{\mathbf{a}'} \exp(\phi^\top(\mathbf{s}, \mathbf{a}')\theta)}$$

or equivalently,

$$\nabla_{\theta} \log \pi(\mathbf{s}, \mathbf{a}; \theta) = \phi(\mathbf{s}, \mathbf{a}) - \mathbb{E}_{\mathbf{a}' \sim \pi(\mathbf{s}, \bullet; \theta)} \phi(\mathbf{s}, \mathbf{a}')$$

Example 2: Continuous Action Space

- Action space: set of real numbers $\mathcal{A} = \mathbb{R}$
- Policy approximator:

$$\pi(\mathbf{s}, \mathbf{a}, \theta) = \frac{1}{\sqrt{2\pi}\sigma(\mathbf{s}; \theta)} \exp\left(-\frac{(\mathbf{a} - \mu(\mathbf{s}; \theta))^2}{2\sigma^2(\mathbf{s}; \theta)}\right),$$

where μ and σ are mean and deviation function approximators

- Linear function approximator with feature vectors $\phi_\mu(\mathbf{s})$ and $\phi_\sigma(\mathbf{s})$
 - $\mu(\mathbf{s}; \theta) = \phi_\mu^\top(\mathbf{s})\theta_\mu$ and $\sigma(\mathbf{s}; \theta) = \phi_\sigma^\top(\mathbf{s})\theta_\sigma$
 - $\nabla_{\theta_\mu} \log \pi(\mathbf{s}, \mathbf{a}, \theta) = \frac{\mathbf{a} - \mu(\mathbf{s}; \theta)}{\sigma^2(\mathbf{s}; \theta)} \phi_\mu(\mathbf{s})$
 - $\nabla_{\theta_\sigma} \log \pi(\mathbf{s}, \mathbf{a}, \theta) = \frac{(\mathbf{a} - \mu(\mathbf{s}; \theta))^2 - \sigma^2(\mathbf{s}; \theta)}{\sigma^2(\mathbf{s}; \theta)} \phi_\sigma(\mathbf{s})$

Example 3: Bernoulli, Logistic Example

- Actions space: binary, $\{0, 1\}$
- Policy approximator:

$$\pi(1, \mathbf{s}; \theta) = 1 - \pi(0, \mathbf{s}; \theta) = \mathbf{p}(\mathbf{s}; \theta)$$

where $\mathbf{p}(\mathbf{s}; \theta)$ is a function approximator

- Linear function approximator with feature vectors $\phi(\mathbf{s})$
 - Logistic function $\sigma(\mathbf{x}) = [\mathbf{1} + \exp(-\mathbf{x})]^{-1}$
 - For exponential soft-max policy $\mathbf{p}(\mathbf{s}; \theta) = \sigma(\phi^\top(\mathbf{s})\theta)$
 - $\nabla_\theta \log(\pi(\mathbf{s}, \mathbf{a}; \theta)) = (\mathbf{a} - \sigma(\phi^\top(\mathbf{s})\theta))\phi(\mathbf{s})$

Lecture Outline

1. Policy Gradient Method Introduction
2. Policy Gradient Theorem
- 3. REINFORCE and Actor Critic Algorithms**
4. Advantage Actor-Critic (A2C)

REINFORCE: MC Policy Gradient Algorithm

- To maximize $J(\theta)$, we apply (stochastic) gradient ascent algorithm

$$\theta_{t+1} = \theta_t + \alpha_t \nabla_{\theta} J(\theta_t)$$

- According to the policy gradient theorem,

$$\nabla_{\theta} J(\theta) = \sum_{\mathbf{s}, \mathbf{a}} \mu^{\pi(\cdot; \theta)}(\mathbf{s}, \mathbf{a}) \nabla_{\theta} \log(\pi(\mathbf{s}, \mathbf{a}; \theta)) Q^{\pi(\cdot; \theta)}(\mathbf{s}, \mathbf{a})$$

- Focus on the average reward setting
- μ^{π} (stationary state-action distribution) is unknown: use empirical state-action distribution $\{(\mathbf{S}_t, \mathbf{A}_t)\}_t$ as an approx
- Q^{π} is unknown: use empirical return $\mathbf{G}_t = \sum_{j=t}^T \mathbf{R}_j$ as an approx

REINFORCE: Pseudocode

- **Initialization:** θ arbitrary
- **For each** episode $(\mathbf{s}_0, \mathbf{a}_0, \mathbf{r}_0, \dots, \mathbf{s}_T, \mathbf{a}_T, \mathbf{r}_T)$ generated using policy $\pi(\bullet; \theta)$

For $t = 0, 1, 2, \dots, T$ **do**:

$$\theta \leftarrow \theta + \alpha \nabla_{\theta} \log(\pi(\mathbf{s}_t, \mathbf{a}_t; \theta)) \mathbf{G}_t$$

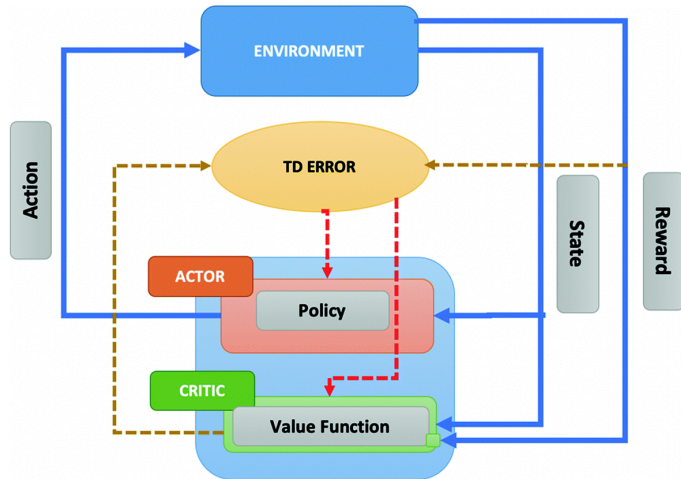
end for

return θ

Actor-Critic Algorithm

- MC policy gradient algorithm may have a large **variance**
 - Return involves many state transitions, many actions and many rewards
- Solution sought by using **actor-critic algorithms**
- Actor-critic algorithms combine **policy gradient** with **value function estimation**

Actor-Critic Algorithm (Cont'd)



- **Critic** uses function approximator to learn value function
- **Actor** uses policy approximator to learn optimal policy

From https://link.springer.com/chapter/10.1007/978-981-13-8285-7_11

Actor-Critic Control

- **Critic:** estimates $Q^{\pi(\cdot;\theta)}(\mathbf{s}, \mathbf{a})$ by a function approximator $\hat{Q}(\mathbf{s}, \mathbf{a}; \omega)$
 - The critic performs **policy evaluation**
 - Standard methods can be applied: MC, TD(0), TD(λ), gradient-based methods
- **Actor:** updates policy parameter θ
 - The actor performs control using approximate policy gradient

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{(\mathbf{s}, \mathbf{a}) \sim \mu} \nabla_{\theta} \log(\pi(\mathbf{s}, \mathbf{a}; \theta)) \hat{Q}(\mathbf{s}, \mathbf{a}; \omega)$$

- Parameter update
 - Average reward setting

$$\theta \leftarrow \theta + \alpha \nabla_{\theta} \log(\pi(\mathbf{S}_t, \mathbf{A}_t; \theta)) \hat{Q}(\mathbf{S}_t, \mathbf{A}_t; \omega)$$

- Discounted reward setting

$$\theta \leftarrow \theta + \alpha \gamma^t \nabla_{\theta} \log(\pi(\mathbf{S}_t, \mathbf{A}_t; \theta)) \hat{Q}(\mathbf{S}_t, \mathbf{A}_t; \omega)$$

Example: Actor-Critic with Linear Value Function

- Linear value function approximator

$$\hat{Q}(s, a; \omega) = \phi^\top(s, a)\omega$$

- Focus on the discounted reward setting
- **Critic**: updates ω by linear TD(0)

$$\omega_{t+1} = \omega_t + \eta \phi(S_t, A_t)(R_t + \gamma \phi^\top(S_{t+1}, A_{t+1})\omega_t - \phi^\top(S_t, A_t)\omega_t)$$

Pseudocode

- **Initialization:** s, θ, ω

- **For each** episode:

Initialize $t = 0$

Sample action a from $\pi(\bullet, s; \theta)$

Repeat until s is terminal

Receive reward r and next state s'

Sample action a' from $\pi(\bullet, s; \theta)$

$$\theta \leftarrow \theta + \alpha \gamma^t \log(\pi(s, a; \theta)) \phi^\top(s, a) \omega$$

$$\omega \leftarrow \omega + \eta \phi(s, a) [r + \gamma \phi^\top(s', a') \omega - \phi^\top(s, a) \omega]$$

$a \leftarrow a'$ and $s \leftarrow s'$

$$t \leftarrow t + 1$$

Bias-Variance Tradeoff

- **REINFORCE** uses Return G_t , an unbiased estimate of $Q^{\pi(\cdot;\theta)}(s, a)$
- **Actor-critic** uses $\hat{Q}(s, a; \omega)$, a biased estimate of $Q^{\pi(\cdot;\theta)}(s, a)$
- REINFORCE gradient has **high variance** and **zero bias**
- Actor-critic gradient has **low variance** and **some bias**
- Similar to Pros & Cons of MC vs TD (Lecture 4, p13)
- Perhaps surprisingly, actor-critic gradient can be **unbiased** under certain conditions (see next slide)

Compatible Function Approximation Theorem

Theorem

Assume the following two conditions:

(C1) *Compatibility of value function approximator and the policy*

$$\nabla_{\omega} \hat{Q}(\mathbf{s}, \mathbf{a}; \omega) = \nabla_{\theta} \log \pi(\mathbf{s}, \mathbf{a}; \theta)$$

(C2) *Value function approximator minimizes the mean-squared error:*

$$\mathbb{E}_{(\mathbf{s}, \mathbf{a}) \sim \mu^{\pi(\cdot; \theta)}} \left[Q^{\pi(\cdot; \omega)}(\mathbf{s}, \mathbf{a}) - \hat{Q}(\mathbf{s}, \mathbf{a}; \omega) \right]^2$$

Then the gradient is unbiased

$$\nabla_{\theta} \mathbf{J}(\theta) = \mathbb{E}_{(\mathbf{s}, \mathbf{a}) \sim \mu^{\pi(\cdot; \theta)}} \nabla_{\theta} (\log \pi(\mathbf{s}, \mathbf{a}; \theta)) \hat{Q}(\mathbf{s}, \mathbf{a}; \omega)$$

Compatible Linear Function Approximation

- Consider the soft-max policy, for a given state-action feature vector $\phi(\mathbf{s}, \mathbf{a})$:

$$\pi(\mathbf{s}, \mathbf{a}; \theta) = \frac{\exp(\phi^\top(\mathbf{s}, \mathbf{a})\theta)}{\sum_{\mathbf{a}'} \exp(\phi^\top(\mathbf{s}, \mathbf{a}')\theta)}$$

- Compatibility condition requires that

$$\nabla_{\omega} \hat{Q}(\mathbf{s}, \mathbf{a}'; \omega) = \nabla_{\theta} \log(\pi(\mathbf{s}, \mathbf{a}; \theta)) = \underbrace{\phi(\mathbf{s}, \mathbf{a}) - \sum_{\mathbf{a}'} \phi(\mathbf{s}, \mathbf{a}')\pi(\mathbf{s}, \mathbf{a}'; \theta)}_{\text{centered state-action features vectors}}$$

which leads to a linear approximation for the value function

$$\hat{Q}(\mathbf{s}, \mathbf{a}'; \omega) = \left[\phi(\mathbf{s}, \mathbf{a}) - \sum_{\mathbf{a}'} \phi(\mathbf{s}, \mathbf{a}')\pi(\mathbf{s}, \mathbf{a}'; \theta) \right]^\top \omega$$

Convergence Theorem

Theorem

Assume

- $\pi(\bullet; \theta)$ and $\hat{Q}(\bullet; \omega)$ are differentiable functions
- Compatibility assumption holds
- The Hessian matrix $\nabla_{\theta}^2 \pi(\mathbf{s}, \mathbf{a}; \theta)$ are uniformly bounded away from infinity
- Step sizes are such that $\sum_t \alpha_t = \infty$ and $\sum_t \alpha_t^2 < \infty$
- At each step, ω_t is chosen to be the solution of

$$\mathbb{E}_{(\mathbf{s}, \mathbf{a}) \sim \mu} \pi(\mathbf{s}, \mathbf{a}; \theta_t) [Q^{\pi(\bullet; \theta_t)}(\mathbf{s}, \mathbf{a}) - \hat{Q}(\mathbf{s}, \mathbf{a}; \omega)] \nabla_{\omega} \hat{Q}(\mathbf{s}, \mathbf{a}; \omega) = 0$$

Then $\{\theta_t\}_t$ are convergent in the sense that $\lim_{t \rightarrow \infty} \|\nabla_{\theta} J(\theta_t)\| \rightarrow 0$.

Separation of Timescales

- The last condition defines ω_t as a solution of a fixed-point equation which has the policy's parameter vector θ_t as a parameter
- In practice, we update ω_t using stochastic gradient descent algorithm. SGD would update ω_t in a similar manner with a larger step size than α_t . It ensures ω converges faster than θ , thus closer to the solution of the fixed-point equation at each time
- This can be seen as a **separation of timescales**:
 - Critic updates the value function approximator at a **faster** timescale trying to evaluate the current policy chosen by the actor
 - Actor varies the policy's parameter more **slowly** to allow the critic to evaluate the current policy
- Similar assumptions are imposed in gradient Q-learning algorithms [Maei et al., 2010]

Lecture Outline

1. Policy Gradient Method Introduction
2. Policy Gradient Theorem
3. REINFORCE and Actor Critic Algorithms
- 4. Advantage Actor-Critic (A2C)**

Variance Reduction Using a Baseline

- Recall that policy parameter update

$$\theta \leftarrow \theta + \alpha \gamma^t \nabla_{\theta} \log(\pi(\mathbf{S}_t, \mathbf{A}_t; \theta)) \hat{Q}(\mathbf{S}_t, \mathbf{A}_t; \omega)$$

- For any θ , when $\mathbf{A}_t \sim \pi(\mathbf{S}_t, \bullet, \theta)$

$$\mathbb{E}[\nabla_{\theta} \log(\pi(\mathbf{S}_t, \mathbf{A}_t, \theta)) | \mathbf{S}_t] = \mathbf{0}$$

- For any baseline function $B(\mathbf{s})$, consider the update

$$\theta \leftarrow \theta + \alpha \gamma^t \nabla_{\theta} \log(\pi(\mathbf{S}_t, \mathbf{A}_t; \theta)) [\hat{Q}(\mathbf{S}_t, \mathbf{A}_t; \omega) - B(\mathbf{S}_t)]$$

- The **mean** of gradient is the same without baseline
- However, the **variance** of the gradient would be smaller with a properly chosen B

Variance Reduction Using a Baseline (Cont'd)

- Consider the baseline that minimizes the variance of the gradient
- For any random variable \mathbf{Z} , the mean $\mathbb{E}\mathbf{Z}$ minimizes $\arg \min_{\mathbf{z}} \mathbb{E}(\mathbf{Z} - \mathbf{z})^2$
- To minimize variance of the gradient $\nabla_{\theta} \log(\pi(\mathbf{S}_t, \mathbf{A}_t; \theta))[\hat{Q}(\mathbf{S}_t, \mathbf{A}_t; \omega) - B(\mathbf{S}_t)]$, the baseline is set to the conditional mean of Q-function given the state
- i.e., $B(\mathbf{s}) = \sum_{\mathbf{a}} \pi(\mathbf{s}, \mathbf{a}; \theta) \hat{Q}(\mathbf{s}, \mathbf{a}; \omega)$, e.g., the estimated state-value
- Similar ideas have been employed in gradient-based algorithms in HW1

Policy Gradient Using Advantage Function

- Advantage function: $A^{\pi(\cdot;\theta)}(\mathbf{s}, \mathbf{a}) = Q^{\pi(\cdot;\theta)}(\mathbf{s}, \mathbf{a}) - V^{\pi(\cdot;\theta)}(\mathbf{s})$
- Policy gradient based on advantage function

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{(\mathbf{s}, \mathbf{a}) \sim \mu^{\pi(\cdot;\theta)}} \nabla_{\theta} \log(\pi(\mathbf{s}, \mathbf{a}; \theta)) A^{\pi(\cdot;\theta)}(\mathbf{s}, \mathbf{a})$$

- The advantage function reduces the variance of policy gradient

An Approach for Estimating Advantage Function

- The critic may compute estimators of both value functions

$$\hat{Q}(\mathbf{s}, \mathbf{a}; \omega) \text{ for } Q^{\pi(\cdot; \theta)}(\mathbf{s}, \mathbf{a})$$

and

$$\hat{V}(\mathbf{s}; \omega) \text{ for } V^{\pi(\cdot; \theta)}(\mathbf{s})$$

which can be done by standard methods such as TD learning

- The estimator of the advantage function

$$\hat{A}(\mathbf{s}, \mathbf{a}; \omega) = \hat{Q}(\mathbf{s}, \mathbf{a}; \omega) - \hat{V}(\mathbf{s}; \omega)$$

Another Approach

- $r + \gamma V^{\pi(\cdot;\theta)}(s') - V^{\pi(\cdot;\theta)}(s)$ is **unbiased** to $A^{\pi(\cdot;\theta)}(s, a)$

$$\begin{aligned} & \mathbb{E}[r + \gamma V^{\pi(\cdot;\theta)}(s') - V^{\pi(\cdot;\theta)}(s) | a, s] \\ &= \mathbb{E}[r + \gamma V^{\pi(\cdot;\theta)}(s') - Q^{\pi(\cdot;\theta)}(s, a) + Q^{\pi(\cdot;\theta)}(s, a) - V^{\pi(\cdot;\theta)}(s) | a, s] \\ &= Q^{\pi(\cdot;\theta)}(s, a) - V^{\pi(\cdot;\theta)}(s) = A^{\pi(\cdot;\theta)}(s, a) \end{aligned}$$

- As such,

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{(s,a) \sim \mu^{\pi(\cdot;\theta)}} \nabla_{\theta} \log(\pi(s, a; \theta)) [r + \gamma V^{\pi(\cdot;\theta)}(s') - V^{\pi(\cdot;\theta)}(s)]$$

- No need to estimate the advantage. It suffices to estimate the state-value and use the estimator to compute the policy gradient

Critic Policy Evaluation Methods

- When specialized to linear methods $\hat{V}(\mathbf{s}; \boldsymbol{\omega}) = \boldsymbol{\phi}^\top(\mathbf{s})\boldsymbol{\omega}$, the critic can use different targets to evaluate

$$\boldsymbol{\omega}_{t+1} \leftarrow \boldsymbol{\omega}_t + \eta_t[\mathbf{v}_t - \boldsymbol{\phi}^\top(\mathbf{S}_t)\boldsymbol{\omega}_t]\boldsymbol{\phi}(\mathbf{S}_t)$$

- The target is defined differently for different methods
 - MC: $\mathbf{v}_t = G_t$
 - TD: $\mathbf{v}_t = R_t + \gamma \hat{V}(\mathbf{S}_{t+1})$
 - TD(λ): $\mathbf{v}_t = G_t^\lambda$

Actor Policy Gradient Methods

- The policy gradient

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{(\mathbf{s}, \mathbf{a}) \sim \mu^{\pi(\cdot; \theta)}} \nabla_{\theta} \log(\pi(\mathbf{s}, \mathbf{a}; \theta)) \mathbf{A}^{\pi(\cdot; \theta)}(\mathbf{s}, \mathbf{a})$$

- Gradient-based method

$$\theta \leftarrow \theta + \alpha \gamma^t \nabla_{\theta} \log(\pi(\mathbf{S}_t, \mathbf{A}_t; \theta)) \hat{\mathbf{A}}(\mathbf{S}_t, \mathbf{A}_t; \omega)$$

- Examples:

- MC: $\hat{\mathbf{A}}(\mathbf{S}_t, \mathbf{A}_t; \omega) = \mathbf{G}_t - \hat{\mathbf{V}}(\mathbf{S}_t; \omega)$
- TD: $\hat{\mathbf{A}}(\mathbf{S}_t, \mathbf{A}_t; \omega) = \mathbf{R}_t + \gamma \hat{\mathbf{V}}(\mathbf{S}_{t+1}; \omega) - \hat{\mathbf{V}}(\mathbf{S}_t; \omega)$

Summary


Policy Function Approximation

		No	Yes
Value Function Approximation	No	Value-based (tabular)	REINFORCE
	Yes	Value-based	Actor-Critic

- **Value-based**
 - Tabular (Lectures 3 & 4)
 - Function approx (Lectures 5 & 7)
- **REINFORCE**
 - No value function
 - Learn policy
- **Actor-critic**
 - Learn value
 - Learn policy
- **Advantage actor-critic**
 - Variance reduction

Seminar Exercise

- Solution to HW7 (Deadline: Wed 12pm)
- Implementation of DQN on LunarLander



Taken from <https://shiva-verma.medium.com/solving-lunar-lander-openaigym-reinforcement-learning-785675066197>

References I

- Hamid Reza Maei, Csaba Szepesvári, Shalabh Bhatnagar, and Richard S Sutton. Toward off-policy learning control with function approximation. In *ICML*, 2010.
- Richard S Sutton, David McAllester, Satinder Singh, and Yishay Mansour. Policy gradient methods for reinforcement learning with function approximation. *Advances in neural information processing systems*, 12, 1999.

Questions

Appendix: Proof of Policy Gradient Theorem

- We focus on the discounted reward setting. Proofs in the average reward setting can be found in Sutton et al. [1999]
- Basic identities

$$(A) \quad V^\pi(\mathbf{s}) = \sum_{\mathbf{a}} \pi(\mathbf{a}|\mathbf{s}) Q^\pi(\mathbf{s}, \mathbf{a})$$

$$(B) \quad Q^\pi(\mathbf{s}, \mathbf{a}) = R_{\mathbf{s}}^{\mathbf{a}} + \gamma \sum_{\mathbf{s}'} P_{\mathbf{s}, \mathbf{s}'}^{\mathbf{a}} V^\pi(\mathbf{s}')$$

$$(C) \quad \nabla_{\theta} V^\pi(\mathbf{s}) = \sum_{\mathbf{a}} [\nabla_{\theta} \pi(\mathbf{a}|\mathbf{s})] Q^\pi(\mathbf{s}, \mathbf{a}) + \sum_{\mathbf{a}} \pi(\mathbf{a}|\mathbf{s}) [\nabla_{\theta} Q^\pi(\mathbf{s}, \mathbf{a})]$$

$$(D) \quad \nabla_{\theta} Q^\pi(\mathbf{s}, \mathbf{a}) = \gamma \sum_{\mathbf{s}'} P_{\mathbf{s}, \mathbf{s}'}^{\mathbf{a}} \nabla_{\theta} V^\pi(\mathbf{s}')$$

Appendix: Proof (Cont'd)

$$\begin{aligned}
 \nabla_{\theta} \mathbf{V}^{\pi}(\mathbf{s}) &\stackrel{(C)}{=} \sum_{\mathbf{a}} [\nabla_{\theta} \pi(\mathbf{a}|\mathbf{s})] \mathbf{Q}^{\pi}(\mathbf{s}, \mathbf{a}) + \sum_{\mathbf{a}} \pi(\mathbf{a}|\mathbf{s}) [\nabla_{\theta} \mathbf{Q}^{\pi}(\mathbf{s}, \mathbf{a})] \\
 &\stackrel{(D)}{=} \sum_{\mathbf{a}} \pi(\mathbf{a}|\mathbf{s}) [\nabla_{\theta} \log(\pi(\mathbf{a}|\mathbf{s}))] \mathbf{Q}^{\pi}(\mathbf{s}, \mathbf{a}) + \underbrace{\gamma \sum_{\mathbf{a}, \mathbf{s}'} \pi(\mathbf{a}|\mathbf{s}) \mathbf{P}_{\mathbf{s}, \mathbf{s}'}^{\mathbf{a}} \nabla_{\theta} \mathbf{V}^{\pi}(\mathbf{s}')}_{I}
 \end{aligned}$$

Now, consider I . Similarly, we have

$$\begin{aligned}
 I &= \sum_{\mathbf{a}, \mathbf{s}', \mathbf{a}'} \pi(\mathbf{a}|\mathbf{s}) \mathbf{P}_{\mathbf{s}, \mathbf{s}'}^{\mathbf{a}} \pi(\mathbf{a}'|\mathbf{s}') [\nabla_{\theta} \log(\pi(\mathbf{a}'|\mathbf{s}'))] \mathbf{Q}^{\pi}(\mathbf{s}', \mathbf{a}') \\
 &\quad + \gamma \sum_{\mathbf{a}, \mathbf{s}', \mathbf{a}', \mathbf{s}''} \pi(\mathbf{a}|\mathbf{s}) \mathbf{P}_{\mathbf{s}, \mathbf{s}'}^{\mathbf{a}} \pi(\mathbf{a}'|\mathbf{s}') \mathbf{P}_{\mathbf{s}', \mathbf{s}''}^{\mathbf{a}'} \nabla_{\theta} \mathbf{V}^{\pi}(\mathbf{s}'')
 \end{aligned}$$

Appendix: Proof (Cont'd)

Recursively applying the first identity, we obtain

$$\nabla_{\theta} \mathbf{V}^{\pi}(\mathbf{s}) = \mu^{\pi(\bullet; \theta)}(\mathbf{s}', \mathbf{a}'; \mathbf{s}) \nabla_{\theta} \log(\pi(\mathbf{s}', \mathbf{a}')) \mathbf{Q}^{\pi}(\mathbf{s}', \mathbf{a}')$$

where

$$\mu^{\pi(\bullet; \theta)}(\mathbf{s}', \mathbf{a}'; \mathbf{s}) = \sum_{t \geq 0} \gamma^t \pi(\mathbf{s}', \mathbf{a}') \Pr^{\pi(\bullet; \theta)}(\mathbf{S}_t = \mathbf{s}' | \mathbf{S}_0 = \mathbf{s})$$