

ST455: Reinforcement Learning

Lecture 9: Model-Based Reinforcement Learning

Chengchun Shi

Lecture Outline

1. What is Model-Based RL
2. How to implement Model-Based RL
3. Simulation-Based Search
4. Mastering the Game of Go

Lecture Outline

1. What is Model-Based RL
2. How to implement Model-Based RL
3. Simulation-Based Search
4. Mastering the Game of Go

Recap: Planning vs Learning

Two fundamental problems in sequential decision making

- **Planning**

- A model of the environment (e.g., state transition, reward function) is **known**
- The agent performs computations with its model, **without** any external interaction
- a.k.a. deliberation, reasoning, introspection, pondering, thought, search

- **Learning**

- The environment is initially **unknown**
- The agent **interacts** with the environment
- The agent **learns** the optimal policy from experience

RL Algorithms We Have Covered So Far

- **Dynamic Programming** (Lecture 3): learn **value** from **model** (planning)
- **MC, TD** (Lectures 3 - 7): learn **value** from **experience** (learning)
- **Policy Gradient** (Lecture 8): learn **policy** from **experience** (learning)
- Today's lecture: **Model-based** RL
 - learn **model** from experience
 - use both learned model and experience to construct a **value** function or **policy**
 - combine learning with planning

What is a Model?

- A model \mathcal{M} is a **representation** of an MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
- The state space \mathcal{S} and action space \mathcal{A} are usually known to us
- The discounted factor γ is **user-specified**
- Only need to learn the state transition \mathcal{P}

$$\mathcal{P}_{ss'}^a = \Pr(\mathcal{S}_{t+1} = s' | \mathcal{S}_t = s, \mathcal{A}_t = a)$$

and reward function \mathcal{R}

$$\mathcal{R}_s^a = \mathbb{E}(\mathcal{R}_t | \mathcal{S}_t = s, \mathcal{A}_t = a)$$

Model-Free v.s. Model-Based RL

- Model-based RL
 - Learn the model (e.g., reward \mathcal{R}_s^a and transition $\mathcal{P}_{ss'}^a$) from experience
 - **Plan** value or policy from model or **integrate** planning with learning
- Model-free RL
 - **Learn** value or policy **without** learning the reward and transition function
 - Rely on Bellman optimality equation
 - Examples: MC, TD, Policy gradient

Model-Free v.s. Model-Based RL (Cont'd)

- **Pros** of model-based RL
- In some applications, we have a **perfect** model (e.g., Go, chess)
- Can handle **offline** data (more in the next lecture)

- **Pros** of model-free RL
- **Dimensional reduction**
- Easier to learn value than model
- # of parameters of $Q^{\pi^{\text{opt}}}$: $|\mathcal{S}||\mathcal{A}|$
- # of parameters of \mathcal{R}_s^a : $|\mathcal{S}||\mathcal{A}|$
- # of parameters of $\mathcal{P}_{ss'}^a$: $|\mathcal{S}|^2|\mathcal{A}|$

Lecture Outline

1. What is Model-Based RL
- 2. How to implement Model-Based RL**
3. Simulation-Based Search
4. Mastering the Game of Go

How to Implement Model-Based RL

- First, we learn a **model** (reward and state transition functions) based on data
- Next, we can implement **planning** based on the learned model
- Alternatively, we can **integrate planning with learning** (Dyna)

How to Implement Model-Based RL

- First, we learn a **model** (reward and state transition functions) based on data
- Next, we can implement **planning** based on the learned model
- Alternatively, we can **integrate planning with learning** (Dyna)

Model Learning

- **Goal:** estimate \mathcal{R}_s^a and $\mathcal{P}_{ss'}^a$ from experience $\{S_0, A_0, R_0, \dots, S_T\}$
- Using supervised learning

$$\begin{aligned} S_0, A_0 &\rightarrow R_0, S_1 \\ S_1, A_1 &\rightarrow R_1, S_2 \\ &\vdots \\ S_{T-1}, A_{T-1} &\rightarrow R_{T-1}, S_T \end{aligned}$$

- Learning $s, a \rightarrow r$ is a **regression** problem
- Learning $s, a \rightarrow s'$ is a **conditional density estimation** problem
- Loss function: least square/Huber loss, KL divergence
- Compute parameter that minimizes empirical loss

Models for Conditional Density Estimation

- Table lookup model
- Conditional kernel density estimation
- Gaussian process model [Williams and Rasmussen, 2006]
- Deep conditional generative learning¹
 - mixture density network [Rothfuss et al., 2019]
 - normalising flows [Trippe and Turner, 2018]

¹<https://deepgenerativemodels.github.io/notes/index.html>

Table Lookup Model

- Finite MDP model
- Count visits $N(\mathbf{s}, \mathbf{a}) = \sum_{t=0}^{T-1} \mathbb{I}(\mathbf{S}_t = \mathbf{s}, \mathbf{A}_t = \mathbf{a})$ to each state-action pair


$$\hat{\mathcal{P}}_{ss'}^{\mathbf{a}} = \frac{1}{N(\mathbf{s}, \mathbf{a})} \sum_{t=0}^{T-1} \mathbb{I}(\mathbf{S}_t = \mathbf{s}, \mathbf{A}_t = \mathbf{a}, \mathbf{S}_{t+1} = \mathbf{s}')$$
$$\hat{\mathcal{R}}_{\mathbf{s}}^{\mathbf{a}} = \frac{1}{N(\mathbf{s}, \mathbf{a})} \sum_{t=0}^{T-1} \mathbb{I}(\mathbf{S}_t = \mathbf{s}, \mathbf{A}_t = \mathbf{a}) R_t$$

- Alternatively
 - At each time step t , record experience tuple $\langle \mathbf{S}_t, \mathbf{A}_t, R_t, \mathbf{S}_{t+1} \rangle$
 - To sample model, based on a state-action pair (\mathbf{s}, \mathbf{a}) , randomly pick tuple matching $\langle \mathbf{s}, \mathbf{a}, \bullet, \bullet \rangle$

Mixture Density Network

- Learn a generic conditional probability mass/density function of \mathbf{Y} given $\mathbf{X} = \mathbf{x}$, $f(\mathbf{y}|\mathbf{x})$ ($\mathbf{Y} = \mathbf{S}_{t+1}$ and $\mathbf{X} = (\mathbf{S}_t, \mathbf{A}_t)$ in our RL setting)
- Combine **Gaussian mixture model** with **deep neural networks**
- Gaussian mixture model has universal approximation property to approximate any **density** function
- Deep neural networks have universal approximation property to approximate any **mean and variance** functions in Gaussian distribution

What is a Gaussian Mixture Model



Taken from <https://towardsdatascience.com/gaussian-mixture-models-explained-6986aaf5a95>

Gaussian Mixture Model (Cont'd)


Model a probability density function $f(\mathbf{y})$ by

$$f(\mathbf{y}) = \sum_{k=1}^K \omega_k \phi(\mathbf{y}; \mu_k, \sigma_k^2),$$

where $\phi(\bullet; \mu, \sigma^2)$ denotes the probability density function of a Gaussian variable with mean μ and variance σ^2 , and ω_k denotes the probability the variable belongs to the k th cluster

Universal Approximation Property

Gaussian mixture model approximates any probability density function as the number of clusters $K \rightarrow \infty$



Mixture Density Network

- Model a probability density function $f(\mathbf{y})$ by

$$f(\mathbf{y}) = \sum_{k=1}^K \omega_k \phi(\mathbf{y}; \mu_k, \sigma_k^2)$$

- We want to model a **conditional** probability density function $f(\mathbf{y}|\mathbf{x})$
- Can be modelled via a **conditional** Gaussian mixture model

$$f(\mathbf{y}|\mathbf{x}) = \sum_{k=1}^K \omega_k(\mathbf{x}) \phi(\mathbf{y}; \mu_k(\mathbf{x}), \sigma_k^2(\mathbf{x}))$$

- Use **deep neural networks** to parametrize $\omega_k(\bullet)$, $\mu_k(\bullet)$ and $\sigma_k^2(\bullet)$

How to Implement Model-Based RL

- First, we learn a **model** (reward and state transition functions) based on data
- Next, we can implement **planning** based on the learned model
- Alternatively, we can **integrate planning with learning** (Dyna)

Planning with Dynamic Programming

- Give a model $\langle \hat{\mathcal{R}}, \hat{\mathcal{P}} \rangle$
- Use dynamic programming algorithm
 - Policy iteration


$$\pi_0 \xrightarrow{\text{red}} V^{\pi_0} \xrightarrow{\text{blue}} \pi_1 \xrightarrow{\text{red}} V^{\pi_1} \xrightarrow{\text{blue}} \dots \xrightarrow{\text{blue}} \pi^{\text{opt}} \xrightarrow{\text{red}} V^{\pi^{\text{opt}}}$$

- Value iteration


$$V^{\pi_0} \xrightarrow{\text{red}} V^{\pi_1} \xrightarrow{\text{red}} V^{\pi_2} \xrightarrow{\text{red}} \dots \xrightarrow{\text{red}} V^{\pi^{\text{opt}}} \xrightarrow{\text{red}} \pi^{\text{opt}}$$

Difference From Model-Free Methods

Sample updates




TD




Q-Learning

Model-based expected updates



Policy Evaluation



Q-Value Iteration

Planning with Model-Free RL

- A simple but powerful approach to planning
- Use the model only to **generate samples**
- **Sample** experience from model:

$$S' \sim \hat{\mathcal{P}}_{S,\bullet}^A \quad \text{and} \quad R = \hat{\mathcal{R}}_S^A$$

- Apply **model-free** RL to samples
 - MC control
 - SARSA
 - Q-learning
- This is often more **efficient** than dynamic programming-based method

Planning with an Inaccurate Model

- Model-based RL computes π^{opt} with respect to the model $\langle \mathcal{S}, \mathcal{A}, \hat{\mathcal{R}}, \hat{\mathcal{P}}, \gamma \rangle$
- Quality of the estimated policy depends heavily on the accuracy of the model
- When model is inaccurate, planning yields a **suboptimal** policy
- **Solution** 1: when model is wrong, using model-free RL
- **Solution** 2: integrate planning with learning

How to Implement Model-Based RL

- First, we learn a **model** (reward and state transition functions) based on data
- Next, we can implement **planning** based on the learned model
- Alternatively, we can **integrate planning with learning** (Dyna)

Real and Simulated Experience


- We consider two sources of experience
- **Real experience:** Sampled from environment (true MDP)

$$\{S_0, A_0, R_0, \dots, S_T\}$$

- **Simulated experience:** Sampled from model (estimated MDP)

$$S' \sim \hat{\mathcal{P}}_{S,\bullet}^A \quad \text{and} \quad R = \hat{\mathcal{R}}_S^A$$


Dyna




Dyna-Q Algorithm

- Initialize $Q(s, a)$ and $\text{model}(s, a)$ for all s and a
- do forever:
 - (a) $s \leftarrow$ current (non-terminal) state
 - (b) $a \leftarrow \epsilon\text{-greedy}(s, Q)$
 - (c) Execute action a ; observe reward r and next state s'
 - (d) $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$
 - (e) $\text{model}(s, a) \leftarrow (r, s')$
 - (f) Repeat n times:
 - $s \leftarrow$ random previously observed state
 - $a \leftarrow$ random action previously taking in s
 - $(r, s') \sim \text{model}(s, a)$
 - $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$

Dyna-Q on a Simple Maze



WITHOUT PLANNING ($n=0$)



WITH PLANNING ($n=50$)




Figure: Policies found through 2nd episode. The arrows indicate greedy action; if no arrow is shown for a state, then all of its action values were equal.


Dyna-Q⁺

- **Motivation:** models maybe **incorrect**; leads to sub-optimal policies
 - **Limited** sample size for a given state-action pair
 - the environment **changes** and new behavior has not been observed
- **Idea:** encourage **long-untried** actions
 - For each state-action pair, check how many times have **elapsed** since it was last tried
 - Use **bonus reward** in action selection:

$$\mathbf{a} \leftarrow \varepsilon - \text{greedy}(\mathbf{s}, \mathbf{Q} + \kappa\sqrt{\tau}),$$


for some small $\kappa > 0$. $\tau(\mathbf{s}, \mathbf{a})$ denotes the times have elapsed since (\mathbf{s}, \mathbf{a}) was last tried

Dyna-Q with an Inaccurate Model



- The changed environment is **easier**
- The left environment was used for the first 3000 steps
- The right environment was used for the rest

Dyna-Q with an Inaccurate Model (Cont'd)



- The changed environment is **harder**
- The left environment was used for the first 1000 steps
- The right environment was used for the rest

Lecture Outline

1. What is Model-Based RL
2. How to implement Model-Based RL
- 3. Simulation-Based Search**
4. Mastering the Game of Go

Two Ways of Planning

- **Background planning**


- Planning is used well **before** an action is selected
- Need to select actions for each state, not **current** state
- Examples: policy iteration and value iteration in Lecture 3

- **Decision-time planning**

- Planning is started and completed **after** encountering each new state S_t
- As a computation to **determine** A_t
- On the next step planning begins anew with S_{t+1} to produce A_{t+1} , and so on

Simulation-Based Search

- Decision-time planning to select the **best** action
- Build a **search tree** with the current state S_t at the root
- **Simulate** episodes of experience from **now** with the model
- Apply **model-free** RL to simulate trajectories



Simple Monte-Carlo Search

- Given a model \mathcal{M} and a **simulation policy** π
- For each action $\mathbf{a} \in \mathcal{A}$
 1. Simulate K episodes from current state \mathbf{S}_t

$$\{\mathbf{S}_t, \mathbf{a}, \mathbf{R}_t^k, \mathbf{S}_{t+1}^k, \mathbf{A}_{t+1}^k, \mathbf{R}_{t+1}^k, \dots, \mathbf{S}_T^k\}_{k=1}^K \sim \mathcal{M}, \pi$$

2. Evaluate actions by mean return (**Monte-Carlo Evaluation**)

$$Q(\mathbf{S}_t, \mathbf{a}) = \frac{1}{K} \sum_{k=1}^K \mathbf{G}_t \rightarrow Q^\pi(\mathbf{S}_t, \mathbf{a})$$

- Select action with maximum value

$$\mathbf{A}_t = \arg \max_{\mathbf{a} \in \mathcal{A}} Q(\mathbf{S}_t, \mathbf{a})$$

Monte-Carlo Tree Search (Evaluation)

- Given a model \mathcal{M}
- Simulate K episodes from current states S_t using current policy π

$$\left\{ S_t, A_t^k, R_t^k, S_{t+1}^k, A_{t+1}^k, R_{t+1}^k, \dots, S_T^k \right\}_{k=1}^K \sim \mathcal{M}, \pi$$

- Build a search tree containing visited states and actions
- Evaluate states $Q(s, a)$ by mean return of episodes from s, a

$$Q(s, a) = \frac{1}{N(s, a)} \sum_{k=1}^K \sum_{u=t}^T \mathbb{I}(S_u = s, A_u = a) G_u \rightarrow Q^\pi(s, a)$$

- After search is finished, select current action with maximum value in search tree

$$A_t = \arg \max_{a \in \mathcal{A}} Q(S_t, a)$$

Monte-Carlo Tree Search (Simulation)

- In MCTS, the simulation policy (rollout policy) π that simulates data **improves**
- Repeat (each simulation)
 - **Evaluate** states $Q(\mathbf{s}, \mathbf{a})$ by Monte-Carlo evaluation
 - **Improve** simulation policy, e.g., by ϵ -greedy(Q)
 - **Monte-Carlo control** applied to **simulated experience**
- Converges to the optimal search tree, $Q(\mathbf{s}, \mathbf{a}) \rightarrow Q^{\pi^{\text{opt}}}(\mathbf{s}, \mathbf{a})$

Lecture Outline

1. What is Model-Based RL
2. How to implement Model-Based RL
3. Simulation-Based Search
- 4. Mastering the Game of Go**

Case Study: the Game of Go



- Invented in China over 2500 years ago
- The **hardest** classic board game
- Much harder than chess:
 - Go has larger number of legal moves than chess (≈ 250 v.s. ≈ 35)
 - Go involve more moves than chess (≈ 150 v.s. ≈ 80)
 - Traditional game-tree search fails in Go

Rules of Go

- Two players place down white and black stones **alternately**
- Stones are **captured** according to simple rules




Figure: Left: the three white stones are not surrounded because point X is unoccupied. Middle: if black places a stone on X, the three white stones are captured and removed from the board. Right: if white places a stone on point X first, the capture is blocked.

- The game ends when neither player wishes to place another stone
- The player with more **territory** wins the game

Two-Player Zero-Sum Markov Games

- Simplest extension of MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{B}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
- \mathcal{A} and \mathcal{B} are actions spaces of first and second players
- \mathcal{R} is reward function. In Go,
 - $R_t = 0$ for all non-terminal steps
 - $R_T = 1$ if Black wins and -1 otherwise
- Let π and ν be policies of the first and second players
- The state-value function depends on both π and ν

$$V^{\pi, \nu}(s) = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R_t \mid S_0 = s, A_t \sim \pi, B_t \sim \nu \right]$$

Causal Diagram




Figure: Causal diagrams for MDPs, TMDPs and POMDPs. Solid lines represent the causal relationships. Dashed lines indicate the information needed to implement the optimal policy. The parallel sign $||$ indicates that the conditional probability function given parent nodes is equal.

Nash Equilibrium

- At each state s , the two players aim to solve two **minimax** problems

$$\arg \max_{\pi} V^{\pi}(s) = \arg \max_{\pi} \min_{\nu} V^{\pi, \nu}(s)$$

$$\arg \min_{\nu} V^{\nu}(s) = \arg \min_{\nu} \max_{\pi} V^{\pi, \nu}(s)$$

- Under **Markov** and **time-homogeneity** assumptions, there exist stationary policies π^* (ν^*) whose values are no worse (better) than any history dependent policy

$$V^{\pi^*, \nu^*}(s) = \arg \max_{\pi} \min_{\nu} V^{\pi, \nu}(s) = \arg \min_{\nu} \max_{\pi} V^{\pi, \nu}(s)$$

similar to the existence of the optimal stationary policy theorem in Lecture 2

- These policies reach a **Nash equilibrium** [Morgenstern and Von Neumann, 1953], i.e., no player can play better by changing his/her own policy

Prisoner's Dilemma

	Prisoner B stays silent (cooperates)	Prisoner B betrays (defects)
Prisoner A stays silent (cooperates)	Each serves 1 year	Prisoner A: 3 years Prisoner B: goes free
Prisoner A betrays (defects)	Prisoner A: goes free Prisoner B: 3 years	Each serves 2 years

Mutual defection is the only Nash equilibrium

Bellman Optimally Equation

- Bellman optimal equation for the **state value**

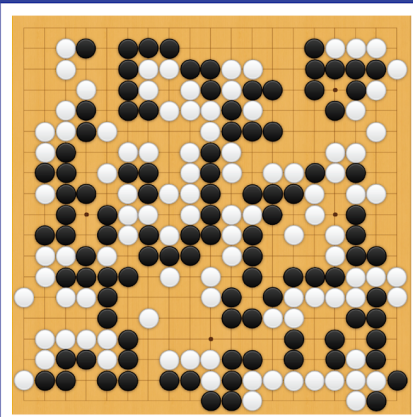
$$V^{\pi^*, \nu^*}(S_t) = \max_a \min_b \mathbb{E} \left[R_t + \gamma V^{\pi^*, \nu^*}(S_{t+1}) \mid S_t, A_t = a, B_t = b \right]$$

- Bellman optimal equation for the **state-action value**

$$Q^{\pi^*, \nu^*}(S_t, A_t, B_t) = \mathbb{E} \left[R_t + \gamma \max_a \min_b Q^{\pi^*, \nu^*}(S_{t+1}, a, b) \mid S_t, A_t = a, B_t = b \right]$$

- The values can be learned similarly to standard TD/Q learning algorithms for MDP [see e.g., Fan et al., 2020]


AlphaGo




THE ULTIMATE GO CHALLENGE
GAME 3 OF 3

27 MAY 2017

● vs ●

 **AlphaGo**
Winner of Match 3


 **Ke Jie**

RESULT B + Res

AlphaGo Pipeline

- Based on a novel version of **Monte-Carlo tree search** (MCTS)
- Combined with a **policy** and a **value function** learned by RL with function approximation provided by deep CNN
- Simulate trajectories and generate the search tree using the **rollout** policy
- **Expand** search tree by selecting unexplored actions according to a **policy network**
- Policy network trained previously via supervised learning to predict moves contained in a database of nearly 30 million human expert moves
- Evaluate state-action value based on simulated returns (MC) and a **value** network
- Value network trained previously via RL

AlphaGo Pipeline (Cont'd)



Input of Neural Networks

Extended Data Table 2 | Input features for neural networks

Feature	# of planes	Description
Stone colour	3	Player stone / opponent stone / empty
Ones	1	A constant plane filled with 1
Turns since	8	How many turns since a move was played
Liberties	8	Number of liberties (empty adjacent points)
Capture size	8	How many opponent stones would be captured
Self-atari size	8	How many of own stones would be captured
Liberties after move	8	Number of liberties after this move is played
Ladder capture	1	Whether a move at this point is a successful ladder capture
Ladder escape	1	Whether a move at this point is a successful ladder escape
Sensibleness	1	Whether a move is legal and does not fill its own eyes
Zeros	1	A constant plane filled with 0
Player color	1	Whether current player is black

Policy Network

- Training the **SL policy network** took approximately 3 weeks using distributed implementation of SGD on 50 processors
- The SL policy network achieved **57%** accuracy; best accuracy achieved by other methods **44%**
- The **RL policy network** is trained on a million games in a single day
- The final RL policy won more than **80%** of games played against the SL policy
- It won **85%** of games played against a Go program using MCTS that simulated 100,000 games per move

Value Network

- The **value network** used **Monte Carlo policy evaluation** based on data obtained from a large number of self-play games played using the RL policy
- To avoid overfitting and instability, and to reduce the strong correlations between positions encountered in self-play, the dataset consists of **30** million positions, each chosen randomly from a unique self-play game
- Training was done using **50** million mini-batches each of **32** positions drawn from this data set
- Training took one week on **50** GPUs

Rollout Policy


- The **rollout policy** was learned prior to play by a simple linear network trained by supervised learning from a corpus of **8** million human moves
- In principle, the SL or RL policy networks could have been used in the rollouts, but the forward propagation through these deep networks took **too much time** for either of them to be used in rollout simulations
- The rollout policy network allowed approximately 1,000 complete game simulations per second to be run on each of the processing threads

Summary

- Model-based/Model free learning
- Integrating planning and learning
- Dyna-Q/Dyna-Q⁺
- Simulation-based search
- Background/Decision-time planning
- Monte Carlo Tree Search
- Two-player zero-sum Markov games
- Nash equilibrium
- AlphaGo


Seminar Exercise

- Solution to HW8 (Deadline: Wed 12:00 pm)



- Advantage Actor-Critic (with deep neural networks) to CartPole

frame: 53, Obs: (0.018, 0.669, 0.286, 0.618)
Action: 1.0, Cumulative Reward: 47.0, Done: 1



- Implementation of Dyna-Q algorithm

References I

- Jianqing Fan, Zhaoran Wang, Yuchen Xie, and Zhuoran Yang. A theoretical analysis of deep q-learning. In *Learning for Dynamics and Control*, pages 486–489. PMLR, 2020.
- Oskar Morgenstern and John Von Neumann. *Theory of games and economic behavior*. Princeton university press, 1953.
- Jonas Rothfuss, Fabio Ferreira, Simon Walther, and Maxim Ulrich. Conditional density estimation with neural networks: Best practices and benchmarks. *arXiv preprint arXiv:1903.00954*, 2019.
- Brian L Trippe and Richard E Turner. Conditional density estimation with bayesian normalising flows. *arXiv preprint arXiv:1802.04908*, 2018.
- Christopher K Williams and Carl Edward Rasmussen. *Gaussian processes for machine learning*, volume 2. MIT press Cambridge, MA, 2006.

Questions